

Lecture 1

- Using the command line
 - The Unix philosophy and `bash`
 - Remote computing with `ssh`
 - Version control with `git`
 - Documents with *L^AT_EX*
 - Automation with `make`

Lecture 2

- The IPython notebook
 - Markdown cells
 - Code cells
 - The display system
 - IPython magic
 - Widgets and the `interact` decorator
 - Interfacing with other languages
- Programming in Python
 - Basic types
 - Basic collections
 - Control flow
 - Functions
 - Classes
 - Modules
 - The standard library
 - PyPI and `pip`
 - Importing other modules
 - Using `conda` virtual environments

Computer lab 1

- Exercise 1: Using `bash` for data scrubbing
- Exercise 2: Using `git` for collaborative editing
- Exercise 3: Generating a report with `make`, *L^AT_EX* and `python`

- Exercise 4: Functions to calculate mean, variance and Pearson correlation coefficient

Lecture 3

- Functional programming
 - Functions are first class objects
 - Pure functions
 - Iterators
 - Generators
 - Anonymous functions with `lambda`
 - Recursion
 - Decorators
 - The `operator` module
 - The `itertools` module
 - The `functools` module
 - The `toolz` module
 - Constructing a lazy data pipeline
- Working with text
 - `string` methods
 - The `string` module
 - The `re` module

Lecture 4

- Obtaining data
 - CSV with `csv`
 - JSON with `json`
 - Web scraping with `scrapy`
 - HDF5 with `pyhdf`
 - Relational databases and SQL with `sqlite3`
 - The `datasets` module
- Scrubbing data
 - Removing comments
 - Filtering rows
 - Filtering columns
 - Fixing inconsistencies
 - Handling missing data
 - Removing unwanted information
 - Derived information
 - Sanity check and visualization

Computer lab 2

- Exercise 1: Moving data from CSV to JSON to HDF to sqlite3
- Exercise 2: Querying a relational database
- Exercise 3: Constructing a functional pipeline using generators
- Exercise 4: [Project Euler puzzle 1](#)

Lecture 5

- Using `numpy`
 - Data types
 - Creating arrays
 - Indexing
 - Broadcasting
 - Outer product
 - Ufuncs
 - Generalized Ufuncs
 - Linear algebra in `numpy`
 - Calculating covariance matrix
 - Solving least squares linear regression
 - I/O in `numpy`
- Using `pandas`
 - Reading and writing data
 - Split-apply-combine
 - Merging and joining
 - Working with time series
- Using `blaze`

Lecture 6

- Graphics in Python
 - Using `matplotlib`
 - Using `seaborn`
 - Using `bokeh`
 - [Using `daft`](#)

Computer lab 3

- Exercise 1: Make a 12 by 12 times table chart without looping
- Exercise 2: Generate a 3X3 magic square
- Exercise 3: Working with some data set in `pandas`
- Exercise 4: Plotting the scatter matrix for the Iris data set in `matplotlib`, `seaborn` and `bokeh`

Lecture 7

- From math to computing
 - Computer representation of numbers
 - Overflow, underflow, catastrophic cancellation
 - Stability
 - Conditioning
 - Direct translation of symbols to code is dangerous
- The purpose of computing is insight not numbers
- Use of the computer in statistics
 - Estimating parameters (point and interval estimates)
 - Estimating functions
 - Approximating functions (especially PDFs)
 - Feature extraction, class discovery and dimension reduction
 - Classification and regression
 - Simulations and computational inference

Lecture 8

- Algorithmic efficiency and big \mathcal{O} notation
- Some data structures
 - Performance of built-in data structures
 - Graphs and trees
- [Algorithmic patterns](#)
 - Divide and conquer
 - Decrease and conquer
 - Greedy algorithms
 - Dynamic programming
 - Hill climbing
 - Reduction and transformation
 - Monte Carlo

Computer lab 4

Exercise 1: [Project Euler puzzle 2](#)

Exercise 2: [Project Euler puzzle 3](#)

Exercise 3: [Project Euler puzzle 4](#)

Exercise 4: [Project Euler puzzle 14](#)

Lecture 9

- Numerical linear algebra
 - Simultaneous linear equations
 - Column space, row space and rank
 - Column space interpretation is most useful
 - Rank, basis, span
 - Norms and distance
 - Trace and determinant
 - Eigenvalues and eigenvectors
 - Inner product
 - Outer product
 - Einstein summation notation
- Matrices as linear transforms
 - Types of matrices
 - Square and non-square
 - Singular
 - Positive definite
 - Idempotent and projections
 - Orthogonal and orthonormal
 - Symmetric
 - Hermitian
 - Transition
 - Matrix geometry illustrated
- Matrix decompositions
 - LU (Gaussian elimination)
 - QR
 - Spectral
 - SVD
 - Cholesky
- Using `scipy.linalg`
- BLAS and LAPACK

Lecture 10

- Projections, ordination, change of coordinates
 - PCA in detail

- PCA with eigendecomposition
- PCA with SVD
- Related methods
 - ICA
 - LSA
 - Factor analysis

Computer lab 5

- Exercise 1: Solving a least squares problem using Cholesky decomposition
- Exercise 2: Implement latent semantic indexing
- Exercise 3: Implement k-means clustering
- Exercise 4: Use latent semantic indexing to reduce a set of documents to 2D then use k-means to cluster them, and finally plot the result

Lecture 11

- Approximating functions
 - Orthogonal function basis
 - Splines
 - Kernel density estimation
- Estimating functions
 - Minimizing least squares
 - Linear regression example and the normal equations
 - Maximum likelihood
 - Logistic regression example
 - Bayesian posterior probability

Lecture 12

- Constrained and unconstrained optimization
- Discrete and continuous optimization
- Root finding and optimization in 1D
 - Taylor series
 - Root finding with Newton methods
 - Root finding with bisection and Brent's method
 - Line search optimization

Computer lab 6

- Exercise 1: Given a sequence of function values, write a program to perform kernel density estimation using several kernels
- Exercise 2: Use line search optimization to solve a 1D logistic regression problem
- Exercise 3: Implement the secant method for finding roots in 1D
- Exercise 4: Implement Newton's method and find an approximate solution to several equations (including ones that diverge)

Lecture 13

- Multivariate optimization
 - GD and SGD
 - Newton's method and IRLS for GLMs
- Other methods
 - Non-gradient based (e.g. Nelder-Mead)
 - Global optimization
 - Discrete optimization
 - Integer programming
 - Combinatorial optimization
- Packages for optimization
 - Using `scipy.optimize`
 - Using `statsmodels`
 - Using `scikit-learn`
 - Others: `cvxopt` and `pyopt`
- General approach to optimization
 - Know the problem
 - Multiple random starts
 - Combining algorithms
 - Graphing progress

Lecture 14

- The EM algorithm (1)
 - Convex and concave functions
 - Jansen's inequality
 - Missing data setup
 - Toy example - coin flipping with 2 biased coins

Computer lab 7

Exercise 1: Write the SGD function to solve a multivariate logistic regression problem using maximum likelihood

Exercise 2: Write the EM algorithm to solve another toy problem

Exercise 3: Playing with `scipy.optimize`

Exercise 4: Playing with scikits-learn

Lecture 15

- The EM algorithm (2)
 - Gaussian mixture model
 - EM for Bayesians - MAP of posterior distribution
 - Other applications of EM
 - EM variants

Lecture 16

- Monte Carlo methods
 - Random number generators
 - Generating random variates from a distribution
 - Quadrature and volume estimation
 - Estimate confidence intervals (bootstrap)
 - Compare competing statistics (statistical simulation - e.g. power)
 - Compare models (cross-validation)
 - Hypothesis testing (permutation-resampling)

Computer lab 8

- Exercise 1: Modify the EM algorithm for GMMs to find the MAP estimate of the posterior distribution
- Exercise 2: Use k-fold cross-validation to evaluate which is the best model for a given data set
- Exercise 3: Estimate the distribution of the slope in a linear regression model by bootstrapping on the residuals
- Exercise 4: Find the type-1 error for $\alpha = 0.05$ by using permutation resampling to correct for multiple testing

Lecture 17

- Conducting a simulation experiment (case study)
- Experimental design
 - Variables to study
 - Levels of variables (factorial, Latin hypercube)
 - Code documentation
 - Recording results
 - Reporting
 - Reproducible analysis with `make` and *L^AT_EX*

Lecture 18

- MCMC (1)
 - Toy problem - rats on drugs
 - Monte Carlo estimation
 - Importance sampling
 - Metropolis-Hasting
 - Gibbs sampling
 - Hamiltonian sampling
 - Assessing convergence
 - Using `pystan`
 - Using `pymc2`
 - Using `emcee`

Computer lab 9

- Exercise 1: Writing a Gibbs sampler for change point detection
- Exercise 2: Using `pystan`
- Exercise 3: Using `pymc2`
- Exercise 4: Using `emcee`

Lecture 19

- MCMC (2)
 - Gaussian mixture model revisited
 - Gibbs sampling
 - Infinite mixture model with the Dirichlet process
 - Simulated tempering
 - MC³

Lecture 20

- Profiling
 - Premature optimization is the root of all evil
 - Using `%time` and `%timeit`
 - Profiling with `%prun`
 - Line profiler
 - Memory profiler
- Code optimization
 - Use appropriate data structure
 - Use appropriate algorithm
 - Use known Python idioms
 - Use optimized modules
 - Caching and memoization
 - Vectorize and broadcast
 - Views
 - Stride tricks

Computer lab 10

- Exercise 1 The label-switching problem
- Exercise 2: Classifying points with the GMM:
- Exercise 3: Profiling source code
- Exercise 4 Optimizing source code

Lecture 21

- JIT compilation with `numba`
- Optimization with `cython`
- Wrapping C code
- Wrapping C++ code
- Wrapping Fortran code

Lecture 22

- [Why modern CPUs are starving and what can be done about it](#)
- Parallel programming patterns
- Amdahl's and Gustafson's laws
- Parallel programming examples
 - JIT compilation with `numba`
 - Toy example - fractals
 - Using `joblib`
 - Using `multiprocessing`
 - Using `IPython.Parallel`
 - Using `MPI4py`

Computer lab 11

Exercise 1: Optimizing EM code with `numba`

Exercise 2: Optimizing EM code with `Cython`

Exercise 3: Parallel processing of embarrassingly parallel code

Exercise 4: Parallel processing of code requiring intr-process communication

Lecture 23

- GPU computing
 - Introduction to CUDA
 - Vanilla matrix multiplication
 - Matrix multiplication with shared memory

- JIT compilation with numba
- Matrix multiplication in OpneCL

Lecture 24

- Map-reduce and Spark
 - Problem - k-mer counting for DNA sequences
 - Small scale map-reduce using Python
 - Using hadoop with mr job
 - Using spark with pyspark
 - Using MLlib for large-scale machine learning

Computer lab 12

Exercise 1: Coding fractals in CUDA

Exercise 2: Something more statistical in CUDA

Exercise 3: Word count in map-reduce

Exercise 4: K-mer count with map reduce with E Coli and human genome

Data sets

- [ecoli genome](#)
- [human genome](#)